

Modeling and Scoring with RevoScaleR

Ali Zaidi

October 11, 2016

Introduction

URL for Today

Please refer to the github repository for course materials
github.com/akzaidi/R-cadence

Agenda

- ▶ We will learn in this tutorial how to train and test models with the RevoScaleR package.
- ▶ Use your knowledge of data manipulation to create **train** and **test** sets.
- ▶ Use the modeling functions in RevoScaleR to train a model.
- ▶ Use the rxPredict function to test/score a model.
- ▶ We will see how you can score models on a variety of data sources.
- ▶ Use a functional methodology, i.e., we will create functions to automate the modeling, validation, and scoring process.

Prerequisites

- ▶ Understanding of `rxDataStep` and `xdfs`
- ▶ Familiarity with `RevoScaleR` modeling and `datastep` functions: `rxLinMod`, `rxGlm`, `rxLogit`, `rxDTree`, `rxDForest`, `rxSplit`, and `rxPredict`
- ▶ Understand how to write functions in R
- ▶ Access to at least one interesting dataset

Typical Lifecycle




Figure 1:

Typical Modeling Lifecycle:

- ▶ Start with a data set
- ▶ Split into a training set and validation set(s)
- ▶ Use the ScaleR modeling functions on the train set to estimate your model
- ▶ Use `rxPredict` to validate/score your results

Mortgage Dataset

- ▶ We will work with a mortgage dataset, which contains mortgage and credit profiles for various mortgage holders

```
mort_path <- paste(rxGetOption("sampleDataDir"), "mortDefault",  
file.copy(mort_path, "mortgage.xdf", overwrite = TRUE))
```

```
## [1] TRUE
```

```
mort_xdf <- RxXdfData("mortgage.xdf")  
rxGetInfo(mort_xdf, getVarInfo = TRUE, numRows = 5)
```

```
## File name: /home/alizaidi/mr4ds/Student-Resources/rmarko
```

```
## Number of observations: 1e+05
```

```
## Number of variables: 6
```

```
## Number of blocks: 10
```

```
## Compression type: zlib
```

```
## Variable information:
```

```
## Var 1: creditScore, Type: integer, Low/High: (470, 925)
```

Transform Default to Categorical

- We might be interested in estimating a classification model for predicting defaults based on credit attributes

```
rxDataStep(inData = mort_xdf,  
           outFile = mort_xdf,  
           overwrite = TRUE,  
           transforms = list(default_flag = factor(ifelse(c  
                                                         )  
                                                         )  
           )  
rxGetInfo(mort_xdf, numRows = 3, getVarInfo = TRUE)
```

```
## File name: /home/alizaidi/mr4ds/Student-Resources/rmarko  
## Number of observations: 1e+05  
## Number of variables: 7  
## Number of blocks: 10  
## Compression type: zlib
```


Modeling

Generating Training and Test Sets

- ▶ The first step to estimating a model is having a tidy training dataset.
- ▶ We will work with the mortgage data and use `rxSplit` to create partitions.
- ▶ `rxSplit` splits an input `.xdf` into multiple `.xdf`s, similar in spirit to the `split` function in base R
- ▶ output is a list
- ▶ First step is to create a split variable
- ▶ We will randomly partition the data into a train and test sample, with 75% in the former, and 25% in the latter

Partition Function

```
create_partition <- function(xdf = mort_xdf,
                             partition_size = 0.75, ...) {
  rxDataStep(inData = xdf,
             outFile = xdf,
             transforms = list(
               trainvalidate = factor(
                 ifelse(rbinom(.rxNumRows,
                               size = 1, prob = splitperc
                             ),
                       "train", "validate")
             )
             ),
             transformObjects = list(splitperc = partition_size,
                                     overwrite = TRUE, ...)

  splitDS <- rxSplit(inData = xdf,
                    #outFilesBase = ,
                    outFileSuffixes = c("train", "validate"),
                    splitByFactor = "trainvalidate",
```

Minimizing IO

Transforms in rxSplit

While the above example does what we want it to do, it's not very efficient. It requires two passes over the data, first to add the `trainvalidate` column, and then another to split it into train and validate sets. We could do all of that in a single step if we pass the transforms directly to `rxSplit`.

```
create_partition <- function(xdf = mort_xdf,
                             partition_size = 0.75, ...) {

  splitDS <- rxSplit(inData = xdf,
                    transforms = list(
                      trainvalidate = factor(
                        ifelse(rbinom(.rxNumRows,
                                     size = 1, prob = sp
                                     "train", "validate")
                      )
                    )
  )
}
```

Generating Training and Test Sets

List of xdfs

- The `create_partition` function will output a list xdfs

```
mort_split <- create_partition(reportProgress = 0)
names(mort_split) <- c("train", "validate")
lapply(mort_split, rxGetInfo)
```

```
## $train
```

```
## File name: /home/alizaidi/mr4ds/Student-Resources/rmarko
```

```
## Number of observations: 75063
```

```
## Number of variables: 8
```

```
## Number of blocks: 10
```

```
## Compression type: zlib
```

```
##
```

```
## $validate
```

```
## File name: /home/alizaidi/mr4ds/Student-Resources/rmarko
```

```
## Number of observations: 24937
```

Build Your Model

Model Formula

- ▶ Once you have a training dataset, the most appropriate next step is to estimate your model
- ▶ RevoScaleR provides a plethora of modeling functions to choose from: decision trees, ensemble trees, linear models, and generalized linear models
- ▶ All take a formula as the first object in their call

```
make_form <- function(xdf = mort_xdf,  
                      resp_var = "default_flag",  
                      vars_to_skip = c("default", "trainval"),  
  
  library(stringr)  
  
  non_incl <- paste(vars_to_skip, collapse = "|")  
  
  x_names <- names(xdf)
```

Build Your Model

Modeling Function

- ▶ Use the `make_form` function inside your favorite `rx` modeling function
- ▶ Default value will be a logistic regression, but can update the `model` parameter to any `rx` modeling function

```
make_form()
```

```
## default_flag ~ creditScore + houseAge + yearsEmploy + co  
##      year  
## <environment: 0x115080db8>
```

```
estimate_model <- function(xdf_data = mort_split[["train"]]  
                           form = make_form(xdf_data),  
                           model = rxLogit, ...) {
```

Build Your Model

Train Your Model with Our Modeling Function

- ▶ Let us now train our logistic regression model for defaults using the `estimate_model` function from the last slide

```
default_model_logit <- estimate_model(mort_split$train,  
                                     reportProgress = 0)  
summary(default_model_logit)
```

```
## Call:  
## model(formula = form, data = xdf_data, reportProgress =  
##  
## Logistic Regression Results for: default_flag ~ creditSc  
##      houseAge + yearsEmploy + ccDebt + year  
## Data: xdf_data (RxxdfData Data Source)  
## File name:  
##      /home/alizaidi/mr4ds/Student-Resources/rmarkdown/mor  
## R plot: xdf_data$default_flag ~ xdf_data$creditSc
```


Building Additional Models

- We can change the parameters of the `estimate_model` function to create a different model relatively quickly

```
default_model_tree <- estimate_model(mort_split$train,
                                     model = rxDTree,
                                     minBucket = 10,
                                     reportProgress = 0)

summary(default_model_tree)
```

##	Length	Class	Mode
## frame	9	data.frame	list
## where	0	-none-	NULL
## call	5	-none-	call
## cptable	30	-none-	numeric
## method	1	-none-	character
## parms	3	-none-	list
## control	9	-none-	list
## splits	65	-none-	numeric

Validation

How Does it Perform on Unseen Data

rxPredict for Logistic Regression

```
## [1] TRUE
```

- ▶ Now that we have built our model, our next step is to see how it performs on data it has yet to see
- ▶ We can use the rxPredict function to score/validate our results

```
default_logit_scored <- rxPredict(default_model_logit,  
                                  mort_split$validate,  
                                  "scored.xdf",  
                                  writeModelVars = TRUE,  
                                  extraVarsToWrite = "default_logit",  
                                  predVarNames = c("pred_logit"))  
  
rxGetInfo(default_logit_scored, numRows = 2)
```

Visualize Model Results

```
rxRoc(actualVarName = "default",  
      predVarNames ="pred_logit_default",  
      data = default_logit_scored)
```

##	threshold	predVarName	sensitivity	specificity
## 1	0.00	pred_logit_default	1.000000000	0.0000000
## 2	0.01	pred_logit_default	0.886178862	0.9407995
## 3	0.02	pred_logit_default	0.780487805	0.9648988
## 4	0.03	pred_logit_default	0.699186992	0.9751753
## 5	0.04	pred_logit_default	0.650406504	0.9804143
## 6	0.05	pred_logit_default	0.609756098	0.9844040
## 7	0.06	pred_logit_default	0.593495935	0.9871847
## 8	0.07	pred_logit_default	0.536585366	0.9889175
## 9	0.08	pred_logit_default	0.471544715	0.9902877
## 10	0.09	pred_logit_default	0.455284553	0.9914967
## 11	0.10	pred_logit_default	0.430894309	0.9925848
## 12	0.11	pred_logit_default	0.430894309	0.9935117
## 13	0.12	pred_logit_default	0.414634146	0.9941162

Testing a Second Model

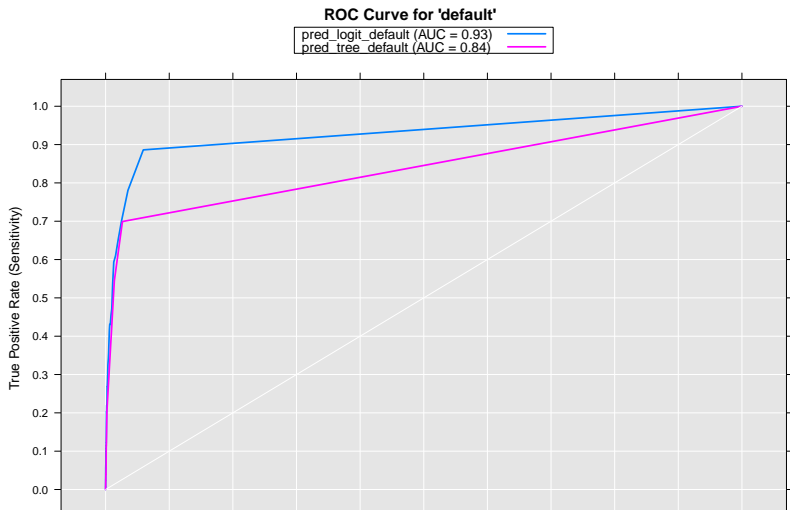
rxPredict for Decision Tree

- ▶ We saw how easy it was to train on different in the previous sections
- ▶ Similarly simple to test different models

[illegible]

Visualize Multiple ROCs

```
rxRocCurve("default",  
           c("pred_logit_default", "pred_tree_default"),  
           data = default_tree_scored)
```



Lab - Estimate Other Models Using the Functions Above

Ensemble Tree Algorithms

Two of the most predictive algorithms in the RevoScaleR package are the `rxBTrees` and `rxDForest` algorithms, for gradient boosted decision trees and random forests, respectively.

Use the above functions and estimate a model for each of those algorithms, and add them to the `default_tree_scored` dataset to visualize ROC and AUC metrics.

```
## Starter code
```

```
default_model_forest <- estimate_model(mort_split$train,  
                                       model = rxDForest,  
                                       nTree = 100,  
                                       importance = TRUE,  
                                       reportProgress = 0)  
  
default_forest_scored <- rxPredict(default_model_forest,  
                                   mort_split$validate,  
                                   "scored.xdf",
```


More Advanced Topics

Scoring on Non-XDF Data Sources

Using a CSV as a Data Source

- ▶ The previous slides focused on using xdf data sources
- ▶ Most of the rx functions will work on non-xdf data sources
- ▶ For training, which is often an iterative process, it is recommended to use xdfs
- ▶ For scoring/testing, which requires just one pass through the data, feel free to use raw data!

```
csv_path <- paste(rxGetOption("sampleDataDir"),  
                  "mortDefaultSmall2009.csv",  
                  sep = "/")  
file.copy(csv_path, "mortDefaultSmall2009.csv", overwrite =
```

```
## [1] TRUE
```

```
mort_csv <- RxTextData("mortDefaultSmall2009.csv")
```

Regression Tree

- ▶ For a slightly different model, we will estimate a regression tree.
- ▶ Just change the parameters in the `estimate_model` function

```
tree_model_ccdebt <- estimate_model(xdf_data = mort_split$train,
                                   form = make_form(mort_split$test,
                                                    "ccDebt",
                                                    vars_t = vars(ccdebt)),
                                   model = rxDTree)

# plot(RevoTreeView::createTreeView(tree_model_ccdebt))
```

Test on CSV

```
if (file.exists("mort2009predictions.xdf")) file.remove("mo
```

```
## [1] TRUE
```

```
rxPredict(tree_model_ccdebt,  
          data = mort_csv,  
          outData = "mort2009predictions.xdf",  
          writeModelVars = TRUE)
```

```
mort_2009_pred <- RxXdfData("mort2009predictions.xdf")  
rxGetInfo(mort_2009_pred, numRows = 1)
```

```
## File name: /home/alizaidi/mr4ds/Student-Resources/rmarko
```

```
## Number of observations: 10000
```

```
## Number of variables: 7
```

```
## Number of blocks: 1
```

```
## Compression type: zlib
```

```
## Data (1 row starting with row 1):
```

Multiclass Classification

Convert Year to Factor

- ▶ We have seen how to estimate a binary classification model and a regression tree
- ▶ How would we estimate a multiclass classification model?
- ▶ Let's try to predict mortgage origination based on other variables
- ▶ Use `rxFactors` to convert *year* to a *factor* variable

```
mort_xdf_factor <- rxFactors(inData = mort_xdf,  
                             factorInfo = c("year"),  
                             outFile = "mort_year.xdf",  
                             overwrite = TRUE)
```

Convert Year to Factor

```
rxGetInfo(mort_xdf_factor, getVarInfo = TRUE, numRows = 4)

## File name: /home/alizaidi/mr4ds/Student-Resources/rmarko
## Number of observations: 1e+05
## Number of variables: 7
## Number of blocks: 10
## Compression type: zlib
## Variable information:
## Var 1: creditScore, Type: integer, Low/High: (470, 925)
## Var 2: houseAge, Type: integer, Low/High: (0, 40)
## Var 3: yearsEmploy, Type: integer, Low/High: (0, 14)
## Var 4: ccDebt, Type: integer, Low/High: (0, 14094)
## Var 5: year
##           10 factor levels: 2000 2001 2002 2003 2004 2005 2
## Var 6: default, Type: integer, Low/High: (0, 1)
## Var 7: default_flag
##           2 factor levels: current default
## Data (4 rows starting with row 1):
```


Predict Multiclass Classification

- Score the results

```
multiclass_preds <- rxPredict(tree_multiclass_year,  
                               data = mort_xdf_factor,  
                               writeModelVars = TRUE,  
                               outData = "multi.xdf",  
                               overwrite = TRUE)
```

Predict Multiclass Classification

- ▶ View the results
- ▶ Predicted/scored column for each level of the response
- ▶ Sum up to one

```
rxGetInfo(multiclass_preds, numRows = 3)
```

Conclusion

Thanks for Attending!

- ▶ Any questions?
- ▶ Try different models!
- ▶ Try modeling with `rxDForest`, `rxBTrees`: have significantly higher predictive accuracy, somewhat less interpretability